

Machine Learning

Homework 5

Not collected, not graded.

1 AdaBoost

1. What is a weak classifier?
2. What is a strong classifier?
3. How does AdaBoost select the weak classifiers?
4. What is the importance of modifying the sample weights between weak classifier selection?

2 K -means clustering

We are given points $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^D$ and an integer $K > 1$, and our goal is to minimize the within cluster variance:

$$J(\{\boldsymbol{\mu}_k\}, \{l_n\}) = \sum_{n=1}^N |\mathbf{x}_n - \boldsymbol{\mu}_{l_n}|^2$$

where $\boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_K \in \mathbb{R}^D$ are the K cluster centers, and $l_1, \dots, l_N \in \{1, \dots, K\}$ are the cluster assignments. You have seen the popular Lloyd's algorithm in class, to find an approximate solution by optimizing $\{\boldsymbol{\mu}_k\}$ and $\{l_n\}$, alternately.

1. Show that Lloyd's algorithm is always guaranteed to converge (=stop) in a finite number of steps. For simplicity assume that $\forall n \in \{1, \dots, N\}: \forall j \neq k: |\mathbf{x}_n - \boldsymbol{\mu}_k|^2 \neq |\mathbf{x}_n - \boldsymbol{\mu}_j|^2$, i.e. there will never be a tie to be broken when assigning clusters. *Hint*: think about how many possible cluster assignments there are, and how the objective function $J(\{\boldsymbol{\mu}_k\}, \{l_n\})$ changes at each step.
2. Does this mean that Lloyd's algorithm always converges to the global optimum?
3. We have assumed that the number of clusters, K , is being revealed to us by some benevolent oracle (in other words: picking the right K is a difficult problem). Explain how the exact minimum of the K -means objective function behaves (on any dataset) as we increase K from 1 to N .
4. If we were to run K -means on a given dataset for all K from 1 to N , and plot the objective function value for each K , could you think of a strategy to deduce the optimal number of clusters from that plot?

3 Expectation–Maximization for Gaussian mixture models

A Gaussian mixture model is a family of distributions of the following form:

$$p(\mathbf{x}) := \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

where the mixture weights satisfy

$$\sum_{k=1}^K \pi_k = 1, \quad \pi_k \geq 0, \quad k \in \{1, \dots, K\}$$

We introduce a latent variable $z_n \in \{1, \dots, K\}$, that encodes the class assignment of sample n . Starting from an initialization for $\boldsymbol{\mu}_k$, $\boldsymbol{\Sigma}_k$, and π_k , $k \in \{1, \dots, K\}$, E-M iterates the following two steps:

1. E-step:

$$\gamma_{n,j} := p(z_n = j \mid \mathbf{x}_n, \{(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, \pi_k)\}_{k=1}^K)$$

i.e., we estimate the “responsibility” $\gamma_{n,j}$ of component j for sample n .

2. M-step:

$$\{(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, \pi_k)\}_{k=1}^K = \arg \max \left\{ \sum_{n=1}^N \sum_{k=1}^K \gamma_{n,k} \ln(\pi_k \mathcal{N}(\mathbf{x}_n \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)) \right\}$$

i.e., we maximize the parameters of each component k considering the n samples weighted by the responsibility $\gamma_{n,k}$.

1. Derive the M-step updates for $\boldsymbol{\mu}_k$.

2. Derive the M-step updates for π_k . *Hint*: remember the summation constraint on π_k !

3. Consider a simplified Gaussian mixture model, where all components share the same covariance matrix $\boldsymbol{\Sigma}_k = \boldsymbol{\Sigma}$. Derive the update rule for $\boldsymbol{\Sigma}$ in the M-step. This answer can rely on the value $\boldsymbol{\mu}_k$ already estimated at the current M-step.

4. Consider an even simpler Gaussian mixture model, where the shared covariance matrix takes the isotropic form, $\boldsymbol{\Sigma} = \sigma^2 I_D$, for some $\sigma^2 > 0$ *known*. Given the data, we estimate the means $\{\boldsymbol{\mu}_k\}$ of each component and the mixture weights $\{\pi_k\}$ with E-M. Assume that

- the mixture weights $\{\pi_k\}$ are bounded away from zero, i.e. $\exists \varepsilon > 0$ such that $\forall k \in \{1, \dots, K\}: \pi_k \geq \varepsilon$, throughout iterations.
- there will be no ties throughout the iterations: $\forall n \in \{1, \dots, N\}: \forall j \neq k: |\mathbf{x}_n - \boldsymbol{\mu}_k|^2 \neq |\mathbf{x}_n - \boldsymbol{\mu}_j|^2$

Show that as $\sigma^2 \rightarrow 0$, the E-step converges to the update rule for the labels l in Lloyd’s algorithm, i.e., the soft assignment becomes hard. *Hint*: the responsibility is computed like a “soft-max”...

4 Kernel lego

1. Consider the candidate kernel $k(\mathbf{x}, \mathbf{y}) := (c + \mathbf{x}^T \mathbf{y})^2$, $c > 0$.

- Using the lego table seen in class (also: PRML p. 296), show that this is an actual kernel.
- Determine the corresponding mapping $\mathbf{x} \mapsto \phi(\mathbf{x})$.

2. Consider the kernel $k(\mathbf{x}, \mathbf{y}) := (\mathbf{x}^T \mathbf{y})^3$. Find the corresponding mapping $\mathbf{x} \mapsto \phi(\mathbf{x})$.

3. Can you find a general expression for the mappings $\mathbf{x} \mapsto \phi(\mathbf{x})$ corresponding to kernels

$$k(\mathbf{x}, \mathbf{y}) := (\mathbf{x}^T \mathbf{y})^M$$

for $M = 1, 2, \dots$?

4. Sketch a proof for $\exp(\mathbf{x}^T \mathbf{y})$ being a valid kernel. Can you sketch a corresponding (infinite dimensional) mapping?

5. The radial basis function kernel, $k(\mathbf{x}, \mathbf{y}) := \exp(-\gamma \|\mathbf{x} - \mathbf{y}\|^2)$ is very popular. Can you prove it is an actual kernel? *Hint*: All “lego”-rules seen in class are fair game. It may be helpful to expand \exp into a series.

6. What about the more complete version, the “Gaussian” kernel, $k(\mathbf{x}, \mathbf{y}) := \exp(-\frac{1}{2}(\mathbf{x} - \mathbf{y})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \mathbf{y}))$?

5 Gaussian processes for regression

Consider the regression model $t_n = y_n + \varepsilon_n$, where y_n is the model prediction and ε is Gaussian noise with precision β . We want to reproduce and re-derive in detail the results seen in class.

1. Give the likelihood $p(t_n | y_n)$. *Hint*: it's Gaussian...
2. Looking at an entire set of $\mathbf{t} = (t_1, \dots, t_N)^T$ and $\mathbf{y} = (y_1, \dots, y_N)^T$, give their joint likelihood $p(\mathbf{t} | \mathbf{y})$ **as a multivariate pdf**.
Hint: it's still Gaussian. Consider the samples to be independent.
3. The principal idea of the Gaussian process tells us that $p(\mathbf{y}) = \mathcal{N}(\mathbf{y} | 0, \mathbf{K})$, where \mathbf{K} is the kernel-based Gram matrix, with $k_{mn} = k(\mathbf{x}_m, \mathbf{x}_n)$. Compute the marginal $p(\mathbf{t}) := \int p(\mathbf{t} | \mathbf{y}) p(\mathbf{y}) d\mathbf{y}$.
4. Extend the previous result to give the $(N + 1)$ -result for $p(t_{(N+1)}, \mathbf{t})$ seen in class.
5. Based on the joint $p(t_{(N+1)}, \mathbf{t})$, use pages 87-90 of PRML to get the conditional $p(t_{(N+1)} | \mathbf{t})$, which will allow to make predictions for new $\mathbf{x}_{(N+1)}$.